



//table of contents /



COVER ART BY I-HUA CHEN

<u>23</u>

OMNIFACES: MAKING JSF A LOT EASIER

By Anghel Leonard
Solve many day-to-day
JSF problems with
a single, integrated
utility library.

31

JAX-RS 2.0: USING ALL THE GOODNESS

By Abhishek Gupta
Understanding the
client API, filters,
interceptors, and other
useful REST features

41

LONG POLLING WITH ASYNCHRONOUS SERVLETS

By Henry Naftulin
The reliable workhorse
of client/server
communications is
the easy-to-use
fallback when other
methods don't work.

47

PUSHING DATA IN BOTH DIRECTIONS WITH WEBSOCKETS

By Danny Coward
Build on WebSockets'
long-lasting connections to create a simple chat app

<u>03</u>

From the Editor

The Rise and Fall of Languages in 2015

06

Letters to the Editor

Comments, questions, suggestions, and kudos

80

Events

Upcoming Java conferences and events

<u>11</u>

Java Books

Reviews of two Java tutorials

<u>59</u>

Path to Java 9

An Early Look at Java 9 Modules

By Ben Evans

Preparing for modularity—the biggest change by far in the next release of Java

65

JVM Languages

Gosu: A Modern, Down-to-Earth Language for the JVM

By Scott McKinney

A low-ceremony language used in enterprise apps offers an extraordinarily flexible, yet static, type system.

<u>72</u>

Containers

Using Multiple Docker Containers

By Arun Gupta

Assemble a cluster of Docker containers and run a Java EE app—without a lot of housekeeping.

<u>80</u>

Microservices

KumuluzEE: Building Microservices with Java EE

By Tilen Faganel and Matjaz B. Juric Develop self-contained microservices with standard Java EE APIs using the open source KumuluzEE framework.

89

Fix This

By Simon Roberts
Our latest code challenges

40

Java Proposals of Interest

JEP 254: Compact Strings

71

User Groups

Manchester Java Community

93

Contact Us

Have a comment? Suggestion? Want to submit an article proposal? Here's how.







EDITORIAL

Editor in Chief

Andrew Binstock

Managing Editor

Claire Breen

Copy Editors

Karen Perkins, Jim Donahue

- - -

Section Development

Michelle Kovac

Technical Reviewers

Stephen Chin, Reza Rahman

DESIGN

Senior Creative Director

Francisco G Delgadillo

Design Director

Richard Merchán

Senior Designer

Arianna Pucherelli

Designer

Jaime Ferrand

Senior Production Manager

Sheila Brennan

Production Designer

Kathy Cygnarowicz

PUBLISHING

Publisher

Jennifer Hamilton +1.650.506.3794

Associate Publisher and Audience Development Director

Karin Kinnear +1.650.506.1985

Audience Development Manager

Jennifer Kurtz

ADVERTISING SALES

Tom Cometa tom.cometa@oracle.com

Advertising Sales Assistant

Cindy Elhaj +1.626.396.9400 x 201

Mailing-List Rentals

Contact your sales representative.

RESOURCES

Oracle Products

+1.800.367.8674 (US/Canada)

Oracle Services

+1.888.283.0591 (US)

ARTICLE SUBMISSION

If you are interested in submitting an article, please *email* the editors.

SUBSCRIPTION INFORMATION

Subscriptions are complimentary for qualified individuals who complete the subscription form.

MAGAZINE CUSTOMER SERVICE

java@halldata.com **Phone** +1.847.763.9635

PRIVACY

Oracle Publishing allows sharing of its mailing list with selected third parties. If you prefer that your mailing address or email address not be included in this program, contact Customer Service.

Copyright © 2016, Oracle and/or its affiliates. All Rights Reserved. No part of this publication may be reprinted or otherwise reproduced without permission from the editors. JAVA MAGAZINE IS PROVIDED ON AN "AS IS" BASIS. ORACLE EXPRESSLY DISCLAIMS ALL WARRANTIES, WHETHER EXPRESS OR IMPLIED. IN NO EVENT SHALL ORACLE BE LIABLE FOR ANY DAMAGES OF ANY KIND ARISING FROM YOUR USE OF OR RELIANCE ON ANY INFORMATION PROVIDED HEREIN. Opinions expressed by authors, editors, and interviewees—even if they are Oracle employees—do not reflect the views of Oracle. The information is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described for Oracle's products remains at the sole discretion of Oracle. Oracle and Java are registered trademarks of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.

Java Magazine is published bimonthly at no cost to qualified subscribers by Oracle, 500 Oracle Parkway, MS OPL-3A, Redwood City, CA 94065-1600.









7 Billion Devices Run Java

ATMs, Smartcards, POS Terminals, Blu-ray Players, Set Top Boxes, Multifunction Printers, PCs, Servers, Routers, Switches, Parking Meters, Smart Meters, Lottery Systems, Airplane Systems, IoT Gateways, Programmable Logic Controllers, Optical Sensors, Wireless M2M Modules, Access Control Systems, Medical Devices, Building Controls, Automobiles...



#1 Development Platform



//microservices /





TILEN FAGANEL AND MATJAZ B. JURIC

KumuluzEE: Building Microservices with Java EE

Develop self-contained microservices with standard Java EE APIs using the open source KumuluzEE framework.

In this article, we explore how to develop microservices with standard Java EE APIs. Microservices have become an important architectural style for developing complex applications composed of small, independent services deployed as individual processes. These services communicate through well-defined APIs. Each microservice is self-contained; is responsible for a single task; is accessible through a lightweight API, such as REST; and has its own lifecycle. Microservices result in a highly decoupled, modular, and reusable architecture for composite applications.

Developing microservices in Java means that not only is the application designed around REST services, but also that each service, the front end, and other components are deployed separately and independently. In other words, instead of packing all the services, business logic, front end, and other modules into a single EAR file and deploying it as a monolith, microservices use individual archives (such as JAR files) for each component. The <u>KumuluzEE framework</u>, as we'll see shortly, automates the tasks related to the configuration and deployment of microservices, specifically on Java EE. [**Note**: KumuluzEE won the Duke's Choice Award at the 2015 JavaOne conference. —*Ed.*]

Microservices with Java EE

Suppose we want to create an online train-reservation ser-

vice people can use to plan their route and buy their tickets. For brevity, we are going to look at a simplified version of this example; it will allow only viewing routes and booking a route via a simple user interface.

Using the traditional approach, we would create a monolithic EAR package that would include our business logic, services, and one or more WAR packages containing the front end. We would deploy the EAR file to an application server, such as Oracle WebLogic Server or GlassFish.

In contrast, using the microservice architecture, we start by separating the responsibilities. Routes, the booking service (reservations), and the UI become three separately developed, configured, and deployed microservices. They are stateless and communicate through REST interfaces—although we could also use SOAP or remote method invocation (RMI). In this way, we create microservices that are domain-defined and follow the "single responsibility" principle with regard to explicit interfaces. We have created a highly modular, decoupled architecture, where each service is responsible for a single, dedicated piece of functionality. Also, each microservice has its own lifecycle. The proposed architecture using microservices is shown in Figure 1.

In the remainder of this article, we use this simplified example. But keep in mind that microservices are best suited for complex applications and systems.







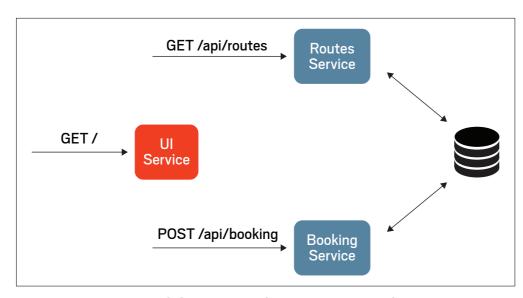


Figure 1. Overview of the proposed microservice architecture

Advantages. Let us now explore the advantages of microservices compared with a monolithic approach. The most obvious advantage is that microservices lead to a more flexible, decoupled architecture. Each microservice can be developed independently, which simplifies lifecycle and change management. For example, if we need to upgrade the booking service, we do not have to redeploy the whole application. Every microservice is a separate project with a separate repository and deployment configuration. We can roll out new features iteratively, thus increasing agility.

If, over time, we would like to upgrade to newer versions of Java, we can upgrade each microservice separately. For complex applications, this is a huge benefit, because it is often difficult to upgrade complex Java EE applications (to a new version of the app server or to new releases of Java EE and Java SE). Even better, microservices allow us to gradually incorporate new technologies.

Last but not least, microservices allow for much better scalability. For example, the route calculation service might get higher traffic than the booking service, because users might use the route calculation more frequently than they book a reservation. With microservices, we can scale the book-

ing service independently of the rest of the application. This approach is a perfect fit for cloud and PaaS (platform as a service) environments, where elastic scalability can be configured easily. Scaling a microservice application in a Docker environment becomes a breeze. In other words, a microservices architecture is a good match for cloud-enabled applications.

To fully exploit these benefits, microservices need to be stateless. Every resource (such as a database, object storage, and so on) that a microservice uses should be configurable separately (usually via a connection string or environment parameters).

Disadvantages. Microservices do not come without drawbacks. Setting up and configuring Java EE projects to accompany this kind of architecture is often not trivial. Actually, it can require a lot of manual configuration and tracking of dependencies, and it can lead to significant operational complexity. Deployment becomes more complex, as does testing.

KumuluzEE

KumuluzEE addresses some of these drawbacks. It automates the tasks related to configuration and deployment of microservices in a seamless way. At its most basic, KumuluzEE collects each microservice together with the Java EE APIs' runtime into a simple, standalone package (JAR file), providing a minimal footprint by including only those APIs (and their runtimes) that are actually used. This enables developers to build standalone, stateless, self-contained microservices and package them in an efficient way without the overhead of including the entire Java EE application server in each microservice. Microservices created this way can be executed directly from the JRE with a minimal footprint and quick startup and shutdown times.

Getting Started

We'll now create our microservices using KumuluzEE, standard Java EE, and Maven. First, we create three Maven







projects—each containing its own microservice. For brevity's sake, we will create them in the same repository. We will create a routes project for the route calculation service, a bookings project for the booking service, and a ui project for the front end. We will also add a module that will hold our JPA entities (models) and a module that will hold our utils, because they will be shared with our three microservices. Note: In real—world projects, it is generally recommended to use a separate repository for each microservice so that they are treated as separate entities and have separate versioning and revision history as well as separate deployment.

We will also create the topmost pom.xml file. You can look at the project structure and the pom.xml file by downloading the source code for this project from <u>GitHub</u>.

To create our microservice projects, we will use Maven archetype generation (mvn -B archetype:generate) for all modules. This will generate the three projects we require.

Adding KumuluzEE

Now we need to add the appropriate dependencies to the KumuluzEE framework. KumuluzEE is completely modular, which means that apart from the core functionality each Java EE component is packaged as a separate module and must be included explicitly as a dependency in order to be used. KumuluzEE will automatically detect which modules are included in the classpath and properly configure them.

We include KumuluzEE by defining the appropriate dependency in Maven. It is recommended that we define a property with the current version of KumuluzEE and use it with every dependency:

Adding the Front-End UI Service

Let's start with the ui microservice and include the core KumuluzEE module for bootstrapping the logic and configuration. It provides the com.kumuluz.ee.EeApplication class with the main method that will bootstrap our app. We will also include the servlet and the HTTP server. We will use a Jetty servlet implementation, which is known for its high performance and small footprint. The ./ui/pom.xml file contains this:

We must also include the maven-dependency-plugin in our pom.xml file, which will copy all our dependencies together with the classes. So we add the following to the ./ui/pom.xml file:







```
</executions>
</plugin>
```

This is the bare minimum required to run a microservice with plain servlets and static files. To try it out, we can add a simple index.xhtml HTML file. KumuluzEE will use the webapp folder at the root of the resource folder to look for configuration data and files. This is where we should put the files.

To try it out, run maven package. Then, you can start the microservice using the following command:

```
$ java -cp ui/target/classes:ui/target/dependency/* \
    com.kumuluz.ee.EeApplication
```

Go to http://localhost:8080/ in your browser and you should see the content of the index.xhtml HTML file. Of course, you can use JSP, use a servlet, or add JSF support and use it for the front end. We will extend the front end shortly.

Defining the JPA Module

We follow the procedure described above for each microservice. Because each microservice is its own project, you can customize it as much as you need.

Before we start developing the routes and booking services, we will define the JPA module with entities. To include JPA, we add the EclipseLink JPA implementation. We will also add the database driver. In our case, we will use the PostgreSQL database. However, you are free to use any database. Place the required dependencies in ./models/pom.xml:

Next, we will add the persistence.xml file and entity classes that will be shared with both our microservices even though the microservices will be run separately. This file is the same as what you would develop in a traditional Java EE application. You can look at the provided sample source code to see an example. Notice that the values defined in persistence.xml for the database URL, username, and password can be overwritten by setting the DATABASE_URL, DATABASE_USER, and DATABASE_PASS environment variables, respectively. Doing this is useful for easy configuration in Docker-style environments.

We are now ready to write the entity classes. We use standard JPA, which does not require any modifications for KumuluzEE. The example JPA entity classes for Route and Booking are shown next. The first file, .../models/Route.java, contains the following:







```
@Temporal(TemporalType.TIMESTAMP)
    private Date duration;
    @ManyToMany(mappedBy = "routes")
    private List<Booking> bookings;
    // ... Omitted getters/setters
 The booking file, .../models/Booking.java, contains
this:
@Entity
@NamedQuery(name="Booking.findAll",
            query="SELECT b FROM Booking b")
public class Booking {
    @Id
    @GeneratedValue(strategy =
        GenerationType.IDENTITY)
    private Integer id;
    @Temporal(TemporalType.TIMESTAMP)
    private Date orderDate;
    @ManyToMany
    private List<Route> routes;
    // ... Omitted getters/setters
```

Implementing the Routes and Booking Services

We are now ready to implement the routes and booking services. We will need JAX-RS for the REST interfaces. We will use <u>Jersey</u> as implemented in KumuluzEE. We will also add CDI, because we want to use the EntityManager injection using @PersistenceContext. However, we could just as well do without CDI to keep our microservice even lighter. We add the following dependencies to ./booking/pom.xml:

Now, let's implement the booking REST service. We'll create a booking resource that contains two methods: createBooking(Booking) and getBooking(). Make sure you add the correct dependencies, as described earlier. Also make sure you add the models module to every microservice for access to JPA entities, whose REST service will be used. The code is:

```
@ApplicationPath("/")
public class BookingApplication extends
    javax.ws.rs.core.Application {
}
```

Then in BookingsResource.java we have this:







84

As you can see, with KumuluzEE the microservice implementation does not require any modifications and is exactly the same as with any other Java EE application. You can implement the remaining microservices the same way or look at the provided sample code.

Handling Service Discovery

Let's suppose that we are using JavaServer Faces (JSF) for the front-end UI module. From there, we will be calling our microservices via REST to receive or save the actual data. The main problem with this is that we need to know the exact URLs of the microservices. We could pass the URLs via environment variables. However, that approach would require manual updating every time anything changes. And in the cloud, these changes can be quite frequent.

A better solution would be to dynamically query the address of the requested microservice. We will use <u>Apache ZooKeeper</u> for service discovery, but we could also use any similar tool. In short, ZooKeeper is a centralized service for maintaining configuration information. We will use it to store the

correct URL to each microservice. We add a helper class, ServiceRegistry, in our utils module to handle the dynamic registering, unregistering, and retrieval of endpoints with ZooKeeper:

```
@ApplicationScoped
public class ServiceRegistry {
    private final CuratorFramework zookeeper;
    private final
        ConcurrentHashMap<String, String> zonePaths;
   @Inject
   public ZooKeeperServiceRegistry()
           throws IOException {
       try {
            String zookeeperUri =
                System.getenv("ZOOKEEPER_URI");
            zookeeper = CuratorFrameworkFactory
                .newClient(zookeeperUri,
                     new RetryNTimes(5, 1000));
            zookeeper.start();
            zonePaths = new ConcurrentHashMap<>();
        } catch (IOException) {
   public void registerService(
            String name, String uri) {
       try {
            String node = "/services/" + name;
            if (zookeeper.checkExists()
                  .forPath(node) == null) {
                zookeeper.create()
                    .creatingParentsIfNeeded()
                    .forPath(node);
            String nodePath =
              zookeeper.create().withMode(
```







The ZOOKEEPER_URL environment variable should contain our ZooKeeper URL. We will be starting a ZooKeeper service alongside our microservices using Docker, which is discussed in the next section. Also make sure that the utils package is added as a dependency to every microservice.

We can now inject the service above into our microservices to register its URL upon startup and unregister it upon shutdown. One way to do this is to add a bean to our booking microservice to handle this task dynamically:

```
@ApplicationScoped
public class BookingService {
    @Inject
    ServiceRegistry services;

    private String serviceName = "trains-booking";
    private String endpointURI;

    public BookingService() {
        endpointURI =
            System.getenv("BASE_URI");
    }

    @PostConstruct
    public void registerService() {
```

We use the @PostConstruct, @PreDestroy, and @ApplicationScoped annotations to make sure we register and unregister the service only once per lifecycle. The BASE_URI environment variable should contain the microservice's public URI. We can now create a similar bean for each of our microservices.

Injecting the Microservices at the Front End

We can use the just-described approach for injection on the front end to connect to the REST services. We can inject our ServiceDiscovery bean to our JSF backing bean to retrieve the URL of a microservice we want to invoke, as shown below:







Updating the JSF Front End

Finally, let's update our index.xhtml file in the ui module to use the created bean and to list the available bookings:

```
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE html >
<html xmlns="http://www.w3.org/1999/xhtml"</pre>
      xmlns:f="http://xmlns.jcp.org/jsf/core"
      xmlns:h="http://xmlns.jcp.org/jsf/html">
  <h:head>
    <title>All available bookings</title>
  </h:head>
  <h:body>
    <h:dataTable
       value="#{bookingsBean.allBookings}"
        var="booking">
      <h:column>
        <f:facet name="header">Booking id</f:facet>
        #{booking.id}
      </h:column>
      <h:column>
        <f:facet name="header">Date of order
        </f:facet>
        #{booking.orderDate}
      </h:column>
    </h:dataTable>
  </h:body>
</html>
```

And we're done! We have created three microservices that will run independently of each other. They communicate through REST interfaces and share common resources such as JPA entities. Of course, we could use HTML5 with AngularJS or ReactJS for the front end as well.

Deploying and Running Using Docker

We can now build and run our microservices. To demonstrate this, we will use Docker. First we need to create a Dockerfile, which will contain the steps to create a Docker image that will build each microservice and run it. Because KumuluzEE runs a microservice as a normal, single-process, standalone JAR file with a minimal footprint (and does not require a separate app server instance for each microservice), it fits perfectly with how Docker operates. The Dockerfile is shown next. Its ./routes/Dockerfile file is in the downloaded files.

As we can see, the Dockerfile is pretty simple and straightforward. It installs the required dependencies, builds the microservice, and supplies the command to run it. In our case, we can use the same Dockerfile for every microservice, except for the final command, which specifies which microservice is going to run in a particular Docker container.

Next, we need to build a Docker image for every microservice that we have:







87

```
$ docker build -t trains/ui \
    -f ui/Dockerfile .
$ docker build -t trains/routes \
    -f routes/Dockerfile .
$ docker build -t trains/bookings \
    -f bookings/Dockerfile .
```

Once built, these images can be run anywhere Docker or the Docker API is used, as well as anywhere Docker Swarm, Kubernetes, various PaaS providers, and many such services are used. To test the images, we can run them locally. We also need to start a ZooKeeper instance by using an existing image. Don't forget to pass the required environment variables to our microservices (using the -e switch). The following example uses 172.17.42.1 as the Docker host, because that is the default. Adjust as needed.

```
$ docker run -p 2181:2181 -d fabric8/zookeeper
$ docker run --name ui -p 3000:8080 -d \
    -e BASE_URI='http://172.17.42.1:3000' \
    -e ZOOKEEPER_URI=172.17.42.1:2181 trains/ui
$ docker run --name routes -p 3001:8080 -d \
    -e BASE_URI='http://172.17.42.1:3001' \
    -e ZOOKEEPER_URI=172.17.42.1:2181 trains/routes
$ docker run --name bookings -p 3002:8080 -d \
    -e BASE_URI='http://172.17.42.1:3002' \
    -e ZOOKEEPER_URI=172.17.42.1:2181 trains/bookings
```

We can now browse our microservice application at http://localhost:3000 or browse our REST services for the routes and bookings modules at http://localhost:3001 and http://localhost:3002, respectively.

Conclusion

In this article, we have seen how to develop microservices with standard Java EE using the open source KumuluzEE framework. As demonstrated, the KumuluzEE framework

automates tasks related to the configuration and deployment of microservices. With simple dependency definitions in Maven, it allows us to create self-contained, standalone JAR files, which contain the required execution environment. Therefore, they can be executed within a standard JRE without requiring an application server. This way, microservices with KumuluzEE provide a minimal footprint and are a good fit for executing in cloud, PaaS, and Docker-style environments, which was demonstrated with this example. We also saw how to handle service discovery using ZooKeeper.

KumuluzEE is an enabler for a microservices architecture in Java EE. The major benefit is that it allows us to use standard Java EE APIs to develop microservices. This way we can also port existing applications to microservices. Although at first blush, microservices might look odd to traditional app server developers, many technologists believe microservices represent the future of Java in the cloud. Feel free to download KumuluzEE and try it out. </article>

Tilen Faganel is lead software architect at Sunesis. He is the lead developer of the KumuluzEE framework for Java, which won the 2015 Duke's Choice Award for best Java innovation.

Matjaz B. Juric, PhD, is a Java Champion and Oracle ACE Director who has authored or coauthored more than 15 books on Java, BPM, and SOA and published in several magazines and journals.



Getting Started with KumuluzEE





